

Sumedh Kumar Prasad Heart logistic Regression

Problem Statement: Finding out the different factors causing Atherosclerotic Heart Disease (AHD)

importing libraries

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

importing dataset

```
In [2]: data=pd.read_csv('Heart.csv',index_col=0)
```

```
In [3]: data.head(5)
```

Out[3]:

	Age	Sex	ChestPain	RestBP	Chol	Fbs	RestECG	MaxHR	ExAng	Oldpeak	Slope
1	63	1	typical	145	233	1	2	150	0	2.3	3
2	67	1	asymptomatic	160	286	0	2	108	1	1.5	2
3	67	1	asymptomatic	120	229	0	2	129	1	2.6	2
4	37	1	nonanginal	130	250	0	0	187	0	3.5	3
5	41	0	nontypical	130	204	0	2	172	0	1.4	1

In [4]: `data.info()`

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 303 entries, 1 to 303
Data columns (total 14 columns):
Age          303 non-null int64
Sex          303 non-null int64
ChestPain    303 non-null object
RestBP       303 non-null int64
Chol         303 non-null int64
Fbs          303 non-null int64
RestECG      303 non-null int64
MaxHR        303 non-null int64
ExAng        303 non-null int64
Oldpeak      303 non-null float64
Slope        303 non-null int64
Ca           299 non-null float64
Thal         301 non-null object
AHD          303 non-null object
dtypes: float64(2), int64(9), object(3)
memory usage: 35.5+ KB
```

In [5]: `data.describe()`

Out[5]:

	Age	Sex	RestBP	Chol	Fbs	RestECG	Max
count	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000
mean	54.438944	0.679868	131.689769	246.693069	0.148515	0.990099	149.6072
std	9.038662	0.467299	17.599748	51.776918	0.356198	0.994971	22.87500
min	29.000000	0.000000	94.000000	126.000000	0.000000	0.000000	71.00000
25%	48.000000	0.000000	120.000000	211.000000	0.000000	0.000000	133.5000
50%	56.000000	1.000000	130.000000	241.000000	0.000000	1.000000	153.0000
75%	61.000000	1.000000	140.000000	275.000000	0.000000	2.000000	166.0000
max	77.000000	1.000000	200.000000	564.000000	1.000000	2.000000	202.0000

Dropping the all the NA values from the dataset

In [6]: `data.dropna(axis=0,inplace=True)`

```
In [7]: data.isnull().any()
```

```
Out[7]: Age           False
Sex             False
ChestPain       False
RestBP          False
Chol            False
Fbs            False
RestECG         False
MaxHR           False
ExAng           False
Oldpeak         False
Slope           False
Ca              False
Thal            False
AHD             False
dtype: bool
```

```
In [8]: data.head()
```

```
Out[8]:
```

	Age	Sex	ChestPain	RestBP	Chol	Fbs	RestECG	MaxHR	ExAng	Oldpeak	Slope
1	63	1	typical	145	233	1	2	150	0	2.3	3
2	67	1	asymptomatic	160	286	0	2	108	1	1.5	2
3	67	1	asymptomatic	120	229	0	2	129	1	2.6	2
4	37	1	nonanginal	130	250	0	0	187	0	3.5	3
5	41	0	nontypical	130	204	0	2	172	0	1.4	1

Finding the correlation between all the numerical variables

```
In [9]: data.corr()
```


```
Out[9]:
```

	Age	Sex	RestBP	Chol	Fbs	RestECG	MaxHR	
Age	1.000000	-0.092399	0.290476	0.202644	0.132062	0.149917	-0.394563	0.0
Sex	-0.092399	1.000000	-0.066340	-0.198089	0.038850	0.033897	-0.060496	0.1
RestBP	0.290476	-0.066340	1.000000	0.131536	0.180860	0.149242	-0.049108	0.0
Chol	0.202644	-0.198089	0.131536	1.000000	0.012708	0.165046	-0.000075	0.0
Fbs	0.132062	0.038850	0.180860	0.012708	1.000000	0.068831	-0.007842	-0.0
RestECG	0.149917	0.033897	0.149242	0.165046	0.068831	1.000000	-0.072290	0.0
MaxHR	-0.394563	-0.060496	-0.049108	-0.000075	-0.007842	-0.072290	1.000000	-0.0
ExAng	0.096489	0.143581	0.066691	0.059339	-0.000893	0.081874	-0.384368	1.0
Oldpeak	0.197123	0.106567	0.191243	0.038596	0.008311	0.113726	-0.347640	0.2
Slope	0.159405	0.033345	0.121172	-0.009215	0.047819	0.135141	-0.389307	0.2
Ca	0.362210	0.091925	0.097954	0.115945	0.152086	0.129021	-0.268727	0.1

Heatmap of correlation variables

```
In [11]: sns.heatmap(data.corr(),linecolor="yellow")
```

```
Out[11]: <matplotlib.axes._subplots.AxesSubplot at 0x26b1e4f89b0>
```

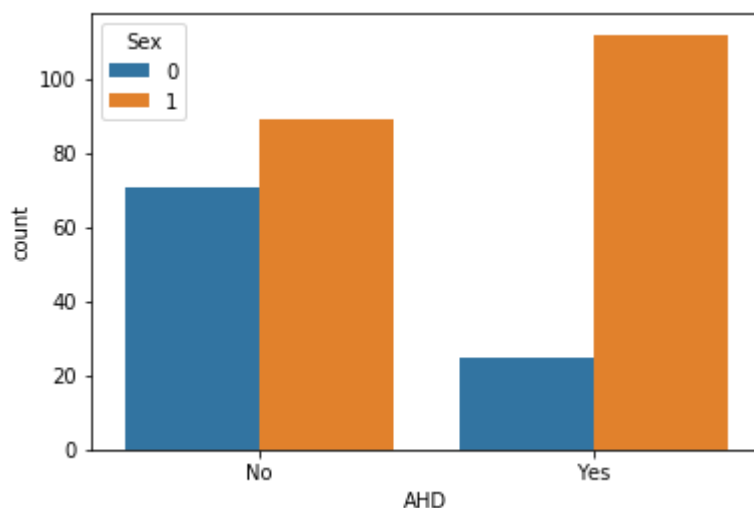


Conclusion from the Heatmap is :

From the corelation we conclude that none of variables are related to each other(No interaction is there between the variables) so we can include all the variables in our dataset for the analysis.

```
In [12]: sns.countplot(x='AHD',hue='Sex', data=data)
```

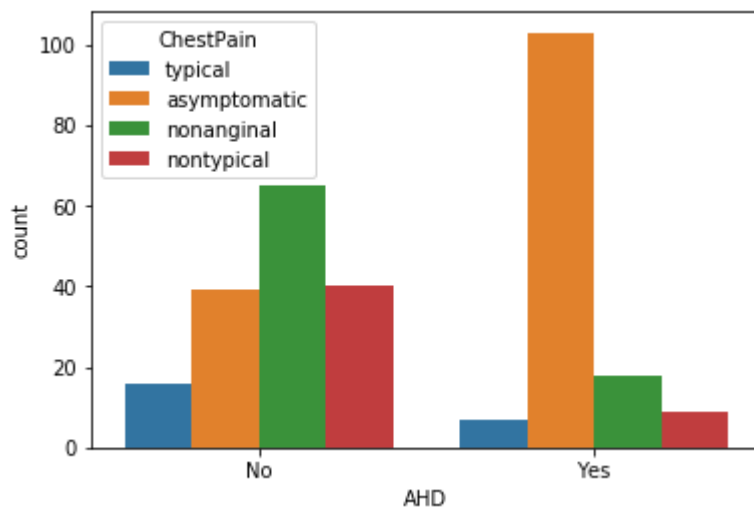
```
Out[12]: <matplotlib.axes._subplots.AxesSubplot at 0x26b1f7f46d8>
```



conclusion : Male has heigher alveolar hydatid disease(AHD) than Female counterpart.

```
In [13]: sns.countplot(x='AHD',hue='ChestPain', data=data)
```

```
Out[13]: <matplotlib.axes._subplots.AxesSubplot at 0x26b1f77c0b8>
```




From the graph it is clear that asymptomatic chest pain causing the Maximum number of AHD attack .

```
In [13]: sns.countplot(x='AHD',hue='Thal', data=data)
```

```
/Users/sanaam/anaconda3/lib/python3.6/site-packages/seaborn/categorical.py:15  
08: FutureWarning: remove_na is deprecated and is a private function. Do not  
use.
```

```
stat_data = remove_na(group_data[hue_mask])
```


```
Out[13]: <matplotlib.axes._subplots.AxesSubplot at 0x11a616a20>
```



From the graph it is clear that Thal having normal not causing the AHD attack most.

```
In [14]: data['Age'].plot.hist()  
plt.xlabel('Age')
```

```
Out[14]: Text(0.5,0,'Age')
```



From the above graph it is clear that Age is approx normally distributed.

Categorical variables (creating dummy variables)

In [14]: `data.columns`

Out[14]: Index(['Age', 'Sex', 'ChestPain', 'RestBP', 'Chol', 'Fbs', 'RestECG', 'MaxHR',
R',
 'ExAng', 'Oldpeak', 'Slope', 'Ca', 'Thal', 'AHD'],
 dtype='object')

In [15]: `data.head(2)`

Out[15]:

	Age	Sex	ChestPain	RestBP	Chol	Fbs	RestECG	MaxHR	ExAng	Oldpeak	Slope
1	63	1	typical	145	233	1	2	150	0	2.3	3
2	67	1	asymptomatic	160	286	0	2	108	1	1.5	2

In [16]: `data1=data.drop(['ChestPain','Thal'],axis=1)`

In [17]: `data1.head(2)`

Out[17]:

	Age	Sex	RestBP	Chol	Fbs	RestECG	MaxHR	ExAng	Oldpeak	Slope	Ca	AHD
1	63	1	145	233	1	2	150	0	2.3	3	0.0	No
2	67	1	160	286	0	2	108	1	1.5	2	3.0	Yes

In [20]: `data=pd.concat([data1,dummy],axis=1)`

In [26]: `data.head(2)`

Out[26]:

	Age	Sex	RestBP	Chol	Fbs	RestECG	MaxHR	ExAng	Oldpeak	Slope	Ca	AHD	Cr
1	63	1	145	233	1	2	150	0	2.3	3	0.0	No	0
2	67	1	160	286	0	2	108	1	1.5	2	3.0	Yes	0

Splitting dataset into dependent (X) & independent variables (Y)

In [27]: `data.columns`

Out[27]: Index(['Age', 'Sex', 'RestBP', 'Chol', 'Fbs', 'RestECG', 'MaxHR', 'ExAng', 'Oldpeak', 'Slope', 'Ca', 'AHD', 'ChestPain_nonanginal', 'ChestPain_nontypical', 'ChestPain_typical', 'Thal_normal', 'Thal_reversable'], dtype='object')

In [28]: `x=data[['Age', 'Sex', 'RestBP', 'Chol', 'Fbs', 'RestECG', 'MaxHR', 'ExAng', 'Oldpeak', 'Slope', 'Ca', 'ChestPain_nonanginal', 'ChestPain_nontypical', 'ChestPain_typical', 'Thal_normal', 'Thal_reversable']]`

`y=data['AHD']`

`y=pd.get_dummies(data['AHD'],drop_first=True)`

In [29]: `x.head(2)`

Out[29]:

	Age	Sex	RestBP	Chol	Fbs	RestECG	MaxHR	ExAng	Oldpeak	Slope	Ca	ChestPa
1	63	1	145	233	1	2	150	0	2.3	3	0.0	0
2	67	1	160	286	0	2	108	1	1.5	2	3.0	0

In [30]: `y.head(2)`

Out[30]:

	Yes
1	0
2	1

Building a Logistic Regression model

Let's start by splitting our data into a training set and test set (there is another test.csv file that you can play around with in case you want to use all this data for training).

Train Test Split of the existing heart dataset.

In [31]: `from sklearn.cross_validation import train_test_split`

/Users/sanaam/anaconda3/lib/python3.6/site-packages/sklearn/cross_validation.py:41: DeprecationWarning: This module was deprecated in version 0.18 in favor of the model_selection module into which all the refactored classes and functions are moved. Also note that the interface of the new CV iterators are different from that of this module. This module will be removed in 0.20.
"This module will be removed in 0.20.", DeprecationWarning)


```
In [32]: x_train, x_test, y_train, y_test = train_test_split( x, y, test_size = 0.3, random_state=101)
```

```
In [33]: y_test[0:5]
```

Out[33]:

	Yes
258	0
275	1
190	1
274	0
136	0

```
In [34]: x_test[0:3].T
```

Out[34]:

	258	275	190
Age	76.0	59.0	69.0
Sex	0.0	1.0	1.0
RestBP	140.0	134.0	140.0
Chol	197.0	204.0	254.0
Fbs	0.0	0.0	0.0
RestECG	1.0	0.0	2.0
MaxHR	116.0	162.0	146.0
ExAng	0.0	0.0	0.0
Oldpeak	1.1	0.8	2.0
Slope	2.0	1.0	2.0
Ca	0.0	2.0	3.0
ChestPain_nonanginal	1.0	0.0	1.0
ChestPain_nontypical	0.0	0.0	0.0
ChestPain_typical	0.0	1.0	0.0
Thal_normal	1.0	1.0	0.0
Thal_reversable	0.0	0.0	1.0

Training and Predicting

```
In [36]: from sklearn.linear_model import LogisticRegression
```

```
logmodel = LogisticRegression()  
logmodel.fit(x_train,y_train)
```

```
/Users/sanaam/anaconda3/lib/python3.6/site-packages/sklearn/utils/validation.py:578: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel().
```

```
y = column_or_1d(y, warn=True)
```

```
Out[36]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,  
    intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,  
    penalty='l2', random_state=None, solver='liblinear', tol=0.0001,  
    verbose=0, warm_start=False)
```

```
In [53]: y_predictions = logmodel.predict(x_test)
```

Finding significant variables

```
In [100]: import statsmodels.api as sm
```

```
In [93]: logit = sm.Logit( y_train, sm.add_constant( x_train ) )
```

```
In [105]: lg = logit.fit()
```

```
Optimization terminated successfully.  
    Current function value: inf  
    Iterations 8
```

```
/Users/sanaam/anaconda3/lib/python3.6/site-packages/statsmodels/discrete/discrete_model.py:1214: RuntimeWarning: overflow encountered in exp  
    return 1/(1+np.exp(-X))
```

```
/Users/sanaam/anaconda3/lib/python3.6/site-packages/statsmodels/discrete/discrete_model.py:1264: RuntimeWarning: divide by zero encountered in log  
    return np.sum(np.log(self.cdf(q*np.dot(X,params))))
```

In [107]: lg.summary()

Out[107]: Logit Regression Results

Dep. Variable:	Yes	No. Observations:	207
Model:	Logit	Df Residuals:	190
Method:	MLE	Df Model:	16
Date:	Sun, 10 Jun 2018	Pseudo R-squ.:	-inf
Time:	15:36:30	Log-Likelihood:	-inf
converged:	True	LL-Null:	-1918.8
		LLR p-value:	1.000

	coef	std err	z	P> z	[0.025	0.975]
const	-1.3550	3.603	-0.376	0.707	-8.417	5.707
Age	-0.0198	0.033	-0.603	0.547	-0.084	0.044
Sex	0.9558	0.631	1.514	0.130	-0.282	2.193
RestBP	0.0279	0.014	1.981	0.048	0.000	0.055
Chol	0.0034	0.005	0.710	0.478	-0.006	0.013
Fbs	-0.8584	0.795	-1.079	0.281	-2.418	0.701
RestECG	0.3986	0.250	1.594	0.111	-0.092	0.889
MaxHR	-0.0351	0.015	-2.398	0.016	-0.064	-0.006
ExAng	0.1698	0.571	0.297	0.766	-0.950	1.290
Oldpeak	0.4138	0.295	1.402	0.161	-0.165	0.992
Slope	0.6533	0.479	1.364	0.173	-0.286	1.592
Ca	1.3000	0.349	3.725	0.000	0.616	1.984
ChestPain_nonanginal	-2.3257	0.665	-3.496	0.000	-3.630	-1.022
ChestPain_nontypical	-1.0752	0.675	-1.592	0.111	-2.399	0.249
ChestPain_typical	-2.5429	0.862	-2.952	0.003	-4.232	-0.854
Thal_normal	0.3314	1.121	0.296	0.767	-1.865	2.528
Thal_reversable	2.1623	1.108	1.952	0.051	-0.009	4.334

```
In [96]: def get_significant_vars( lm ):
          var_p_vals_df = pd.DataFrame( lm.pvalues )
          var_p_vals_df['vars'] = var_p_vals_df.index
          var_p_vals_df.columns = ['pvals', 'vars']
          return list( var_p_vals_df[var_p_vals_df.pvals <= 0.05]['vars'] )
```

```
In [97]: significant_vars = get_significant_vars( lg )
```

```
In [98]: significant_vars
```

```
Out[98]: ['RestBP', 'MaxHR', 'Ca', 'ChestPain_nonanginal', 'ChestPain_typical']
```

Evaluation of the model

We can check Accuracy, precision and recall, using classification report!

Confusion Matrix

```
In [44]: from sklearn.metrics import confusion_matrix
```

```
In [49]: cm=confusion_matrix(y_test,predictions)
```

```
In [50]: cm
```

```
Out[50]: array([[47,  6],
                [ 8, 29]])
```

Graphical representation in terms of confusion matrix


```
In [71]: import seaborn as sn
def draw_cm( actual, predicted ):
    cm = metrics.confusion_matrix( actual, predicted, [1,0] )
    sn.heatmap(cm, annot=True, fmt='.2f', xticklabels = ["Default", "No Default"], yticklabels = ["Default", "No Default"] )
    plt.ylabel('True label')
    plt.xlabel('Predicted label')
    plt.show()
```

```
In [72]: from sklearn import metrics
```

```
In [66]: metrics.confusion_matrix(y_test,y_predictions)
```

```
Out[66]: array([[47,  6],
                [ 8, 29]])
```

```
In [73]: draw_cm( y_test, y_predictions )
```



Note: the model could only predict very few default classes.

Drawing the ROC plot for the Model


```
In [75]: auc_score = metrics.roc_auc_score( y_test, y_predictions )
         round( float( auc_score ), 2 )
```

Out[75]: 0.84

```
In [76]: def draw_roc( actual, probs ):
         fpr, tpr, thresholds = metrics.roc_curve( actual, probs,
                                                    drop_intermediate = False )
         auc_score = metrics.roc_auc_score( actual, probs )
         plt.figure(figsize=(6, 4))
         plt.plot( fpr, tpr, label='ROC curve (area = %0.2f)' % auc_score )
         plt.plot([0, 1], [0, 1], 'k--')
         plt.xlim([0.0, 1.0])
         plt.ylim([0.0, 1.05])
         plt.xlabel('False Positive Rate or [1 - True Negative Rate]')
         plt.ylabel('True Positive Rate')
         plt.title('Receiver operating characteristic example')
         plt.legend(loc="lower right")
         plt.show()

         return fpr, tpr, thresholds
```

```
In [78]: fpr, tpr, thresholds = draw_roc( y_test, y_predictions )
        ## if blue line is below 45 degree it means model is not doing good.
```



```
In [79]: from sklearn.metrics import classification_report
```

```
In [80]: print(classification_report(y_test,predictions))
```

	precision	recall	f1-score	support
0	0.85	0.89	0.87	53
1	0.83	0.78	0.81	37
avg / total	0.84	0.84	0.84	90

```
In [81]: print( 'Total Accuracy : ',np.round( metrics.accuracy_score( y_test, y_predictions), 2 ) )
        print( 'Precision : ',np.round( metrics.precision_score( y_test, y_predictions ), 2 ) )
        print( 'Recall : ',np.round( metrics.recall_score( y_test, y_predictions ), 2 ) )
```

```
cm1 = metrics.confusion_matrix( y_test, y_predictions, [1,0] )
```

```
sensitivity = cm1[0,0]/(cm1[0,0]+cm1[0,1])
print('Sensitivity : ', round( sensitivity, 2 ) )
```

```
specificity = cm1[1,1]/(cm1[1,0]+cm1[1,1])
print('Specificity : ', round( specificity, 2 ) )
```

```
Total Accuracy : 0.84
Precision : 0.83
Recall : 0.78
Sensitivity : 0.78
Specificity : 0.89
```

Total accuracy of my model is 84% .

Conclusion:

Our model predicted with the accuracy of 84% that (AHD) Atherosclerotic Heart Disease is caused by the above given independent variables like(ChestPain RestBP Cholesterol ,Fbs & so on.) because there is no correlation between the any variable. Precision is 83% and Recall is 78% so there is not much difference between both, as both values close to each other and it is greater than 70%. Area under the Curve is the 84% which is above the 45 degree line which shows that Model is good.

Target variable for this Model is Atherosclerotic Heart Disease (AHD)

It is a type of Classification problem within that "Logistics Regression".